For this individual assignment, you will design and build your own LLM-based agent. Your agent can be designed for any purpose you would like, but it must meet the following requirements:

- You must write the code for the agent loop (planning, tool-calling) yourself - you cannot use a framework (e.g. no LangChain, CrewAI etc)
- You MAY use an LLM via API for your agent if you wish, or run your own model
- Your agent must be able to use at least 3 tools (which you can build, or use via API or MCP)
- You must develop a user interface for interacting with your agent
- You must build and execute an approach to evaluating your agent's performance, and demonstrate quantitatively how it performs

The deliverables for this assignment are:

- A 6-minute-max video explaining the agent you built and showing it in action (should explain the goal, tools, design choices and show interface and your evaluation). Can include slides or just verbal + demo. This is NOT a code walkthrough so should be no code shown
- Link to your GitHub repo (be sure to add Akhil and myself if private). Repo should be well organized and include a README with full description of your agent and how to run it